

PersonalizedCancerDiagnosis-Hyperparameter

March 21, 2019

```
In [3]: import pandas as pd
import matplotlib.pyplot as plt
import re
import time
import warnings
import numpy as np
from nltk.corpus import stopwords
from sklearn.decomposition import TruncatedSVD
from sklearn.preprocessing import normalize
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.manifold import TSNE
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics.classification import accuracy_score, log_loss
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import SGDClassifier
from imblearn.over_sampling import SMOTE
from collections import Counter
from scipy.sparse import hstack
from sklearn.multiclass import OneVsRestClassifier
from sklearn.svm import SVC
from sklearn.model_selection import StratifiedKFold
from collections import Counter, defaultdict
from sklearn.calibration import CalibratedClassifierCV
from sklearn.naive_bayes import MultinomialNB
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
import math
from sklearn.metrics import normalized_mutual_info_score
from sklearn.ensemble import RandomForestClassifier
warnings.filterwarnings("ignore")

from mlxtend.classifier import StackingClassifier

from sklearn import model_selection
```

```

from sklearn.linear_model import LogisticRegression

In [4]: train_df = pd.read_pickle('train_df')
test_df = pd.read_pickle('test_df')
cv_df = pd.read_pickle('cv_df')
from scipy import sparse

train_gene_feature_onehotCoding = sparse.load_npz("train_gene_feature_onehotCoding.npz")
train_variation_feature_onehotCoding = sparse.load_npz("train_variation_feature_onehotCoding.npz")
test_gene_feature_onehotCoding = sparse.load_npz("test_gene_feature_onehotCoding.npz")
test_variation_feature_onehotCoding = sparse.load_npz("test_variation_feature_onehotCoding.npz")
cv_gene_feature_onehotCoding = sparse.load_npz("cv_gene_feature_onehotCoding.npz")
cv_variation_feature_onehotCoding = sparse.load_npz("cv_variation_feature_onehotCoding.npz")

In [5]: train_y = np.array(list(train_df['Class']))
test_y = np.array(list(test_df['Class']))
cv_y = np.array(list(cv_df['Class']))
y_train=np.load('y_train.npy')
y_test=np.load('y_test.npy')
y_cv=np.load('y_cv.npy')

In [47]: y_train

Out[47]: array([4, 1, 6, ..., 3, 5, 6], dtype=int64)

In [15]: print('Number of data points in train data:', train_df.shape[0])
print('Number of data points in test data:', test_df.shape[0])
print('Number of data points in cross validation data:', cv_df.shape[0])

Number of data points in train data: 2124
Number of data points in test data: 665
Number of data points in cross validation data: 532

```

0.1 Analysis on Gene Feature

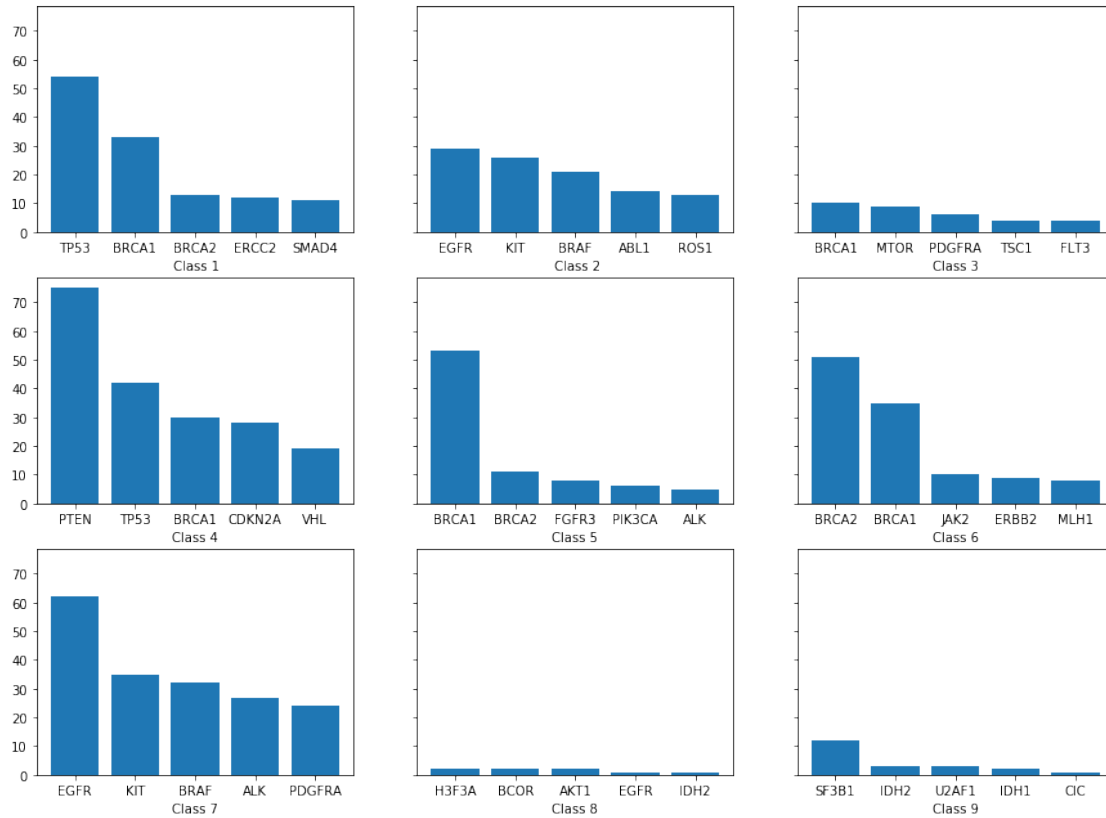
0.2 Gene Frequency corresponding to each class

```

In [139]: fig, axes = plt.subplots(nrows=3, ncols=3, sharey=True, figsize=(15,11))
i=1
p=[' Class 1', ' Class 2', 'Class 3', 'Class 4', 'Class 5', 'Class 6', 'Class 7', 'Class 8']
for row in axes:
    for col in row:
        c={}
        c=train_df[train_df['Class']==i]['Gene'].value_counts()[0:5]

        a=c.to_dict()
        col.bar(list(a.keys()),list(a.values()))
        col.set_xlabel(p[i-1])
        i+=1

```



We can see that BRCA1 and BRCA2 has occurred many times in class 1,3,4,5,6.

0.2.1 Frequency of most occurring Genes corresponding to each class

```
In [5]: train_df['Gene'].value_counts()[0:5]
```

```
Out[5]: BRCA1      161
        TP53       103
        EGFR       98
        PTEN       83
        BRCA2       77
        Name: Gene, dtype: int64
```


These are the most occurring genes

```
In [56]: fig, axes = plt.subplots(nrows=3, ncols=2, sharey=True, figsize=(15,11))
        i=0
        p=[' Class 1',' Class 2','Class 3','Class 4','Class 5','Class 6','Class 7','Class 8',
        gene=['BRCA1','TP53','EGFR','PTEN','BRCA2']
        for row in axes:
            for col in row:
                c={}
                if i<5:
```

```

c=train_df[train_df['Gene']==gene[i]]['Class'].value_counts()
a=c.to_dict()
col.bar(list(a.keys()),list(a.values()))
col.set_xlabel(gene[i])
i+=1
axes[-1, -1].axis('off')
plt.show()

```



1. We can see that BRCA1 which is most occurring gene does not occur even once in class 2/7/8/9.
2. Class 7 contains more number of EGFR gene.
3. Class 9 does not contain any of the most occurring genes.

0.2.2 Analysis on Variation Feature

```

In [8]: counts=train_df['Variation'].value_counts().reset_index()
counts.columns=['Occurrence', 'Count']
count_dict={}
count_dict=counts.Count.value_counts()
total=sum(count_dict.values)
for a, b in count_dict.iteritems():
    print('{}% of variations has {} occurrence'.format(np.round((b/total)*100,2 ),a))

```

98.65% of variations has 1 occurrence
 0.98% of variations has 2 occurrence
 0.1% of variations has 3 occurrence
 0.05% of variations has 60 occurrence
 0.05% of variations has 47 occurrence
 0.05% of variations has 45 occurrence
 0.05% of variations has 19 occurrence
 0.05% of variations has 4 occurrence

We can infer that 98.6% of variations are unique.

0.2.3 Analysis on Text Feature


```
In [113]: train_df.head(5)
```

```
Out[113]:
```

	ID	Gene	Variation	Class	\
	2426	BRCA1	C64G	4	
	636	CDKN1A	Truncating Mutations	1	
	703	ERBB2	E812K	6	
	1455	FGFR2	K659N	7	
	103	MSH6	R976H	1	

		TEXT	word_len	char_len
2426	published analyses effects missense mutations ...		17926	130265
636	introduction loss control mammalian cell cycle...		7234	53047
703	purpose mutations associated resistance kinase...		4961	37512
1455	activating mutations tyrosine kinase domain re...		8077	59908
103	msh6 gene one mismatch repair genes involved h...		13352	100388

```
In [11]: plt.figure(figsize=(12,8))
sns.violinplot(x="Class", y="word_len", data=train_df)
plt.title('Word Length of each class')
plt.show()
```



we can see that most classes have word length between 0 to 10000.

0.2.4 Words with top TFIDF Values

```
In [162]: fig, axes = plt.subplots(nrows=3, ncols=3, sharey=True, figsize=(15,11))
count=1
classes=[' Class 1',' Class 2','Class 3','Class 4','Class 5','Class 6','Class 7','Class 8','Class 9']

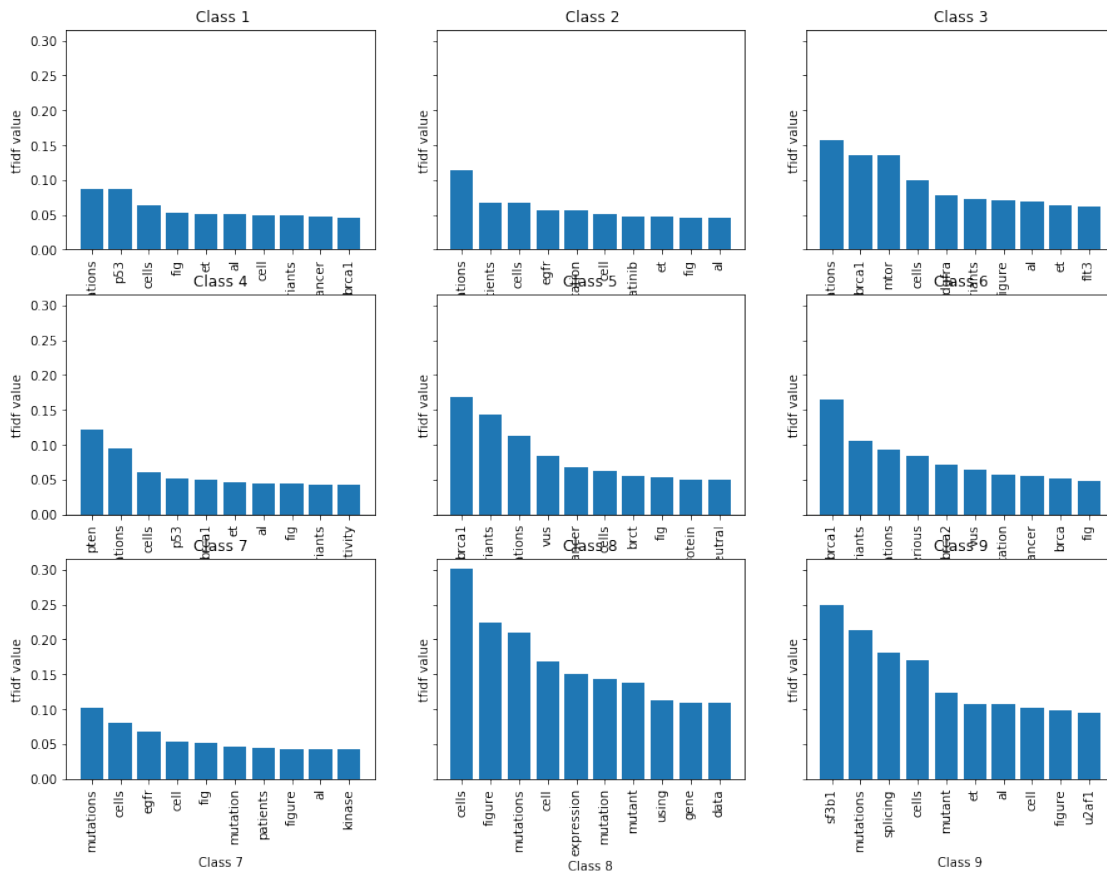
for row in axes:
    for col in row:
        tf_idf_vect = TfidfVectorizer(min_df=10)
        vec1=tf_idf_vect.fit_transform(train_df[train_df['Class']==count]['TEXT'])

        features=pd.DataFrame(tf_idf_vect.get_feature_names())
        mean_tf=np.mean(vec1,axis=0)
        x=np.array(mean_tf)[0].tolist()
        idf_values=sorted(x, reverse=True)
        # taking top 20 features
        important_feat=[]
        important_feat=np.argsort((x))[:, :-1]
        important_feat=important_feat[:10]
        important_feat
        imp_feat=[]
        for index in important_feat:
```

```

        imp_feat.append(features.iloc[index])
x=np.array(imp_feat)
feature=[]
for i in x:
    for j in i:
        feature.append(j)
clas={}
c=0
common={}
for i in feature:
    common[i]=idf_values[c]
    c+=1
col.bar(list(common.keys()),list(common.values()))
col.set_xticklabels(list(common.keys()),rotation=90)
col.set_xlabel(classes[count-1])
col.set_title(classes[count-1])
col.set_ylabel('tfidf value')
count+=1

```



0.3 ML Models

```
In [4]: # This function plots the confusion matrices given y_i, y_i_hat.
def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)
    #C = 9,9 matrix, each cell (i,j) represents number of points of class i are predic

    A = (((C.T)/(C.sum(axis=1))).T)
    #divid each element of the confusion matrix with the sum of elements in that column

    # C = [[1, 2],
    #      [3, 4]]
    # C.T = [[1, 3],
    #        [2, 4]]
    # C.sum(axis = 1) axis=0 corresponds to columns and axis=1 corresponds to rows in
    # C.sum(axis=1) = [[3, 7]]
    # ((C.T)/(C.sum(axis=1))) = [[1/3, 3/7]
    #                             [2/3, 4/7]]

    # ((C.T)/(C.sum(axis=1))).T = [[1/3, 2/3]
    #                               [3/7, 4/7]]
    # sum of row elements = 1

    B = (C/C.sum(axis=0))
    #divid each element of the confusion matrix with the sum of elements in that row
    # C = [[1, 2],
    #      [3, 4]]
    # C.sum(axis = 0) axis=0 corresponds to columns and axis=1 corresponds to rows in
    # C.sum(axis=0) = [[4, 6]]
    # (C/C.sum(axis=0)) = [[1/4, 2/6],
    #                      [3/4, 4/6]]

    labels = [1,2,3,4,5,6,7,8,9]
    # representing A in heatmap format
    print("-"*20, "Confusion matrix", "-"*20)
    plt.figure(figsize=(20,7))
    sns.heatmap(C, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.show()

    print("-"*20, "Precision matrix (Columm Sum=1)", "-"*20)
    plt.figure(figsize=(20,7))
    sns.heatmap(B, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.show()
```



```

# representing B in heatmap format
print("-"*20, "Recall matrix (Row sum=1)", "-"*20)
plt.figure(figsize=(20,7))
sns.heatmap(A, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()

```

```

In [5]: # cls_text is a data frame
# for every row in data fram consider the 'TEXT'
# split the words by space
# make a dict with those words
# increment its count whenever we see that word

```

```

def extract_dictionary_paddle(cls_text):
    dictionary = defaultdict(int)
    for index, row in cls_text.iterrows():
        for word in row['TEXT'].split():
            dictionary[word] +=1
    return dictionary

```

```

In [6]: dict_list = []
# dict_list =[] contains 9 dictoinaries each corresponds to a class
for i in range(1,10):
    cls_text = train_df[train_df['Class']==i]
    # build a word dict based on the words in that class
    dict_list.append(extract_dictionary_paddle(cls_text))
    # append it to dict_list

```

```

# dict_list[i] is build on i'th class text data
# total_dict is buid on whole training text data
total_dict = extract_dictionary_paddle(train_df)

```

```

In [7]: def predict_and_plot_confusion_matrix(train_x, train_y, test_x, test_y, clf):
    clf.fit(train_x, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x, train_y)
    pred_y = sig_clf.predict(test_x)

    # for calculating log_loss we willl provide the array of probabilities belongs to
    print("Log loss :", log_loss(test_y, sig_clf.predict_proba(test_x)))
    # calculating the number of data points that are misclassified
    print("Number of mis-classified points :", np.count_nonzero((pred_y - test_y))/test_y.size)
    plot_confusion_matrix(test_y, pred_y)

```

```

In [8]: def report_log_loss(train_x, train_y, test_x, test_y, clf):
    clf.fit(train_x, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x, train_y)

```

```

sig_clf_probs = sig_clf.predict_proba(test_x)
return log_loss(test_y, sig_clf_probs, eps=1e-15)

```

```

In [9]: # this function will be used just for naive bayes
# for the given indices, we will print the name of the features
# and we will check whether the feature present in the test point text or not
def get_impfeature_names(indices, text, gene, var, no_features):
    gene_count_vec = CountVectorizer()
    var_count_vec = CountVectorizer()
    text_count_vec = CountVectorizer(min_df=3)

    gene_vec = gene_count_vec.fit(train_df['Gene'])
    var_vec = var_count_vec.fit(train_df['Variation'])
    text_vec = text_count_vec.fit(train_df['TEXT'])

    fea1_len = len(gene_vec.get_feature_names())
    fea2_len = len(var_count_vec.get_feature_names())

    word_present = 0
    for i,v in enumerate(indices):
        if (v < fea1_len):
            word = gene_vec.get_feature_names()[v]
            yes_no = True if word == gene else False
            if yes_no:
                word_present += 1
                print(i, "Gene feature [{}]" .format(word))
        elif (v < fea1_len+fea2_len):
            word = var_vec.get_feature_names()[v-(fea1_len)]
            yes_no = True if word == var else False
            if yes_no:
                word_present += 1
                print(i, "variation feature [{}]" .format(word))
        else:
            word = text_vec.get_feature_names()[v-(fea1_len+fea2_len)]
            yes_no = True if word in text.split() else False
            if yes_no:
                word_present += 1
                print(i, "Text feature [{}]" .format(word))

    print("Out of the top ",no_features," features ", word_present, "are present in qu

In [10]: def get_impfeature_names(text, indices, removed_ind = []):
    word_present = 0
    tabulte_list = []
    incresingorder_ind = 0
    for i in indices:
        if i < train_gene_feature_onehotCoding.shape[1]:
            tabulte_list.append([incresingorder_ind, "Gene", "Yes"])

```

```

elif i < 18:
    tabulte_list.append([increasingorder_ind, "Variation", "Yes"])
if ((i > 17) & (i not in removed_ind)) :
    word = train_text_features[i]
    yes_no = True if word in text.split() else False
    if yes_no:
        word_present += 1
    tabulte_list.append([increasingorder_ind, train_text_features[i], yes_no])
increasingorder_ind += 1
print(word_present, "most important features are present in our query point")
print("-"*50)
print("The features that are most important of the ", predicted_cls[0], " class:")
print(tabulate(tabulte_list, headers=["Index", "Feature name", "Present or Not"]))

```

5. Assignments

- Apply All the models with tf-idf features (Replace CountVectorizer with TfidfVectorizer and)
- Instead of using all the words in the dataset, use only the top 1000 words based on tf-idf
- Apply Logistic regression with CountVectorizer Features, including both unigrams and bigrams
- Try any of the feature engineering techniques discussed in the course to reduce the CV and

In [23]: train_df.head()

```

Out[23]:
   ID  Gene  Variation  Class  \
2426 2426  BRCA1      C64G     4
636   636  CDKN1A  Truncating Mutations  1
703   703  ERBB2      E812K     6
1455 1455  FGFR2      K659N     7
103   103  MSH6      R976H     1

                                     TEXT
2426  published analyses effects missense mutations ...
636   introduction loss control mammalian cell cycle...
703   purpose mutations associated resistance kinase...
1455  activating mutations tyrosine kinase domain re...
103   msh6 gene one mismatch repair genes involved h...

```

In [24]: train_variation_feature_onehotCoding

```

Out[24]: <2124x1964 sparse matrix of type '<class 'numpy.int64'>'
         with 2431 stored elements in Compressed Sparse Row format>

```

```

In [224]: # building a CountVectorizer with all the words that occurred minimum 3 times in train
from sklearn.feature_extraction.text import TfidfVectorizer
text_vectorizer = TfidfVectorizer()
train_text_feature = text_vectorizer.fit_transform(train_df['TEXT'])
# getting all the feature names (words)
train_text_features = text_vectorizer.get_feature_names()
test_text_feature = text_vectorizer.transform(test_df['TEXT'])
cv_text_feature = text_vectorizer.transform(cv_df['TEXT'])

```

```
In [13]: print(train_variation_feature_onehotCoding.shape)
         print(train_gene_feature_onehotCoding.shape)
         train_text_feature.shape
```

```
(2124, 1964)
```

```
(2124, 229)
```

```
Out[13]: (2124, 126099)
```

```
In [225]: train_text_feature = normalize(train_text_feature, axis=0)
```

```
# we use the same vectorizer that was trained on train data
```

```
# don't forget to normalize every feature
```

```
test_text_feature = normalize(test_text_feature, axis=0)
```

```
# we use the same vectorizer that was trained on train data
```

```
# don't forget to normalize every feature
```

```
cv_text_feature = normalize(cv_text_feature, axis=0)
```

```
In [29]: train_vec = hstack((train_gene_feature_onehotCoding,train_variation_feature_onehotCoding,train_text_feature_onehotCoding))
         test_vec = hstack((test_gene_feature_onehotCoding,test_variation_feature_onehotCoding,test_text_feature_onehotCoding))
         cv_vec = hstack((cv_gene_feature_onehotCoding,cv_variation_feature_onehotCoding,cv_text_feature_onehotCoding))
```

0.4 Naive Bayes

```
In [52]: alpha = [0.00001, 0.0001, 0.001, 0.1, 1, 10, 100,1000]
```

```
cv_log_error_array = []
```

```
for i in alpha:
```

```
    print("for alpha =", i)
```

```
    clf = MultinomialNB(alpha=i)
```

```
    clf.fit(train_vec, train_y)
```

```
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
```

```
    sig_clf.fit(train_vec, train_y)
```

```
    sig_clf_probs = sig_clf.predict_proba(cv_vec)
```

```
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
```

```
# to avoid rounding error while multiplying probabilities we use log-probability
```

```
    print("Log Loss :",log_loss(cv_y, sig_clf_probs))
```

```
fig, ax = plt.subplots()
```

```
ax.plot(np.log10(alpha), cv_log_error_array,c='g')
```

```
for i, txt in enumerate(np.round(cv_log_error_array,3)):
```

```
    ax.annotate((alpha[i],str(txt)), (np.log10(alpha[i]),cv_log_error_array[i]))
```

```
plt.grid()
```

```
plt.xticks(np.log10(alpha))
```

```
plt.title("Cross Validation Error for each alpha")
```

```
plt.xlabel("Alpha i's")
```

```


plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = MultinomialNB(alpha=alpha[best_alpha])
clf.fit(train_vec, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_vec, train_y)

predict_y = sig_clf.predict_proba(train_vec)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log_loss(train_vec, predict_y))
predict_y = sig_clf.predict_proba(cv_vec)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:",log_loss(cv_vec, predict_y))
predict_y = sig_clf.predict_proba(test_vec)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",log_loss(test_vec, predict_y))

for alpha = 1e-05
Log Loss : 1.2567927668300418
for alpha = 0.0001
Log Loss : 1.2629049300903104
for alpha = 0.001
Log Loss : 1.2686550032314035
for alpha = 0.1
Log Loss : 1.2770262355822313
for alpha = 1
Log Loss : 1.313619241850694
for alpha = 10
Log Loss : 1.3615423725765872
for alpha = 100
Log Loss : 1.2457778208833945
for alpha = 1000
Log Loss : 1.188761554861517

```



For values of best alpha = 1000 The train log loss is: 0.8376778669235427
 For values of best alpha = 1000 The cross validation log loss is: 1.188761554861517
 For values of best alpha = 1000 The test log loss is: 1.1764751736807895

Testing the model with best hyper paramters

```
In [55]: clf = MultinomialNB(alpha=alpha[best_alpha])
         clf.fit(train_vec, train_y)
         sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
         sig_clf.fit(train_vec, train_y)
         sig_clf_probs = sig_clf.predict_proba(cv_vec)
         # to avoid rounding error while multiplying probabillites we use log-probability estim
         print("Log Loss :", log_loss(cv_y, sig_clf_probs))
         print("Number of missclassified point :", np.count_nonzero((sig_clf.predict(cv_vec) - cv_y)))
         plot_confusion_matrix(cv_y, sig_clf.predict(cv_vec.toarray()))
```

Log Loss : 1.188761554861517

Number of missclassified point : 0.39473684210526316


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



0.4.1 Feature Importance, Correctly classified point

```
In [11]: # this function will be used just for naive bayes
# for the given indices, we will print the name of the features
# and we will check whether the feature present in the test point text or not
# feature importance for tfidf features
def get_impfeature_names_tf(indices, text, gene, var, no_features):
    gene_count_vec = CountVectorizer()
    var_count_vec = CountVectorizer()
    text_count_vec = TfidfVectorizer()

    gene_vec = gene_count_vec.fit(train_df['Gene'])
    var_vec = var_count_vec.fit(train_df['Variation'])
    text_vec = text_count_vec.fit(train_df['TEXT'])

    fea1_len = len(gene_vec.get_feature_names())
    fea2_len = len(var_count_vec.get_feature_names())

    word_present = 0
    for i,v in enumerate(indices):
        if (v < fea1_len):
            word = gene_vec.get_feature_names()[v]
            yes_no = True if word == gene else False
            if yes_no:
                word_present += 1
            print(i, "Gene feature [{}] present in test data point [{}]".format(w
        elif (v < fea1_len+fea2_len):
            word = var_vec.get_feature_names()[v-(fea1_len)]
            yes_no = True if word == var else False
            if yes_no:
                word_present += 1
```



```

        print(i, "variation feature [{}]" + "present in test data point [{}]" + "format(w
    else:
        word = text_vec.get_feature_names()[v-(fea1_len+fea2_len)]
        yes_no = True if word in text.split() else False
        if yes_no:
            word_present += 1
        print(i, "Text feature [{}]" + "present in test data point [{}]" + "format(w

    print("Out of the top ", no_features, " features ", word_present, "are present in q

```

In [35]: test_vec=test_vec.tocsr()

```

In [57]: test_point_index = 1
        no_feature = 100
        predicted_cls = sig_clf.predict(test_vec[test_point_index])
        print("Predicted Class :", predicted_cls[0])
        test_vec=test_vec.tocsr()
        print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_vec[test_
        print("Actual Class :", test_y[test_point_index])
        indices = np.argsort(-clf.coef_)[predicted_cls-1][:, :no_feature]
        print("-"*50)
        get_impfeature_names_tf(indices[0], test_df['TEXT'].iloc[test_point_index], test_df['G

```

Predicted Class : 7

Predicted Class Probabilities: [[0.0161 0.0605 0.0191 0.0315 0.033 0.0213 0.8125 0.0038 0.002

Actual Class : 7

```

-----
15 Text feature [cells] present in test data point [True]
16 Text feature [kinase] present in test data point [True]
20 Text feature [cell] present in test data point [True]
21 Text feature [presence] present in test data point [True]
22 Text feature [activated] present in test data point [True]
23 Text feature [contrast] present in test data point [True]
24 Text feature [activation] present in test data point [True]
25 Text feature [phosphorylation] present in test data point [True]
26 Text feature [inhibitor] present in test data point [True]
28 Text feature [also] present in test data point [True]
29 Text feature [shown] present in test data point [True]
30 Text feature [expressing] present in test data point [True]
31 Text feature [10] present in test data point [True]
32 Text feature [signaling] present in test data point [True]
33 Text feature [however] present in test data point [True]
34 Text feature [treated] present in test data point [True]
35 Text feature [independent] present in test data point [True]
37 Text feature [growth] present in test data point [True]
38 Text feature [suggest] present in test data point [True]
39 Text feature [found] present in test data point [True]
40 Text feature [compared] present in test data point [True]

```

41 Text feature [increased] present in test data point [True]
 42 Text feature [similar] present in test data point [True]
 43 Text feature [previously] present in test data point [True]
 45 Text feature [tyrosine] present in test data point [True]
 46 Text feature [inhibition] present in test data point [True]
 47 Text feature [treatment] present in test data point [True]
 48 Text feature [constitutively] present in test data point [True]
 50 Text feature [addition] present in test data point [True]
 52 Text feature [potential] present in test data point [True]
 53 Text feature [well] present in test data point [True]
 54 Text feature [figure] present in test data point [True]
 56 Text feature [various] present in test data point [True]
 57 Text feature [mutations] present in test data point [True]
 58 Text feature [sensitive] present in test data point [True]
 59 Text feature [obtained] present in test data point [True]
 60 Text feature [may] present in test data point [True]
 63 Text feature [showed] present in test data point [True]
 64 Text feature [described] present in test data point [True]
 65 Text feature [3b] present in test data point [True]
 67 Text feature [constitutive] present in test data point [True]
 68 Text feature [inhibitors] present in test data point [True]
 70 Text feature [consistent] present in test data point [True]
 71 Text feature [mutant] present in test data point [True]
 72 Text feature [reported] present in test data point [True]
 73 Text feature [total] present in test data point [True]
 74 Text feature [including] present in test data point [True]
 75 Text feature [using] present in test data point [True]
 76 Text feature [respectively] present in test data point [True]
 77 Text feature [mutation] present in test data point [True]
 81 Text feature [confirmed] present in test data point [True]
 83 Text feature [fig] present in test data point [True]
 84 Text feature [proliferation] present in test data point [True]
 87 Text feature [followed] present in test data point [True]
 88 Text feature [inhibited] present in test data point [True]
 89 Text feature [activating] present in test data point [True]
 90 Text feature [concentrations] present in test data point [True]
 91 Text feature [identified] present in test data point [True]
 92 Text feature [approximately] present in test data point [True]
 95 Text feature [two] present in test data point [True]
 96 Text feature [approved] present in test data point [True]
 97 Text feature [molecular] present in test data point [True]
 98 Text feature [due] present in test data point [True]
 99 Text feature [3a] present in test data point [True]
 Out of the top 100 features 64 are present in query point

0.4.2 Feature Importance, Incorrectly classified point

```
In [68]: test_point_index = 350
         no_feature = 100
         predicted_cls = sig_clf.predict(test_vec[test_point_index])
         print("Predicted Class :", predicted_cls[0])
         print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_vec[test_
         print("Actual Class :", test_y[test_point_index])
         indices = np.argsort(-clf.coef_)[predicted_cls-1][:,no_feature]
         print("-"*50)
         get_impfeature_names_tf(indices[0], test_df['TEXT'].iloc[test_point_index],test_df['G
```

Predicted Class : 1

Predicted Class Probabilities: [[0.5857 0.0275 0.0171 0.2585 0.0318 0.024 0.0492 0.0036 0.002

Actual Class : 4

```
-----
13 Text feature [protein] present in test data point [True]
14 Text feature [type] present in test data point [True]
15 Text feature [wild] present in test data point [True]
16 Text feature [one] present in test data point [True]
19 Text feature [dna] present in test data point [True]
20 Text feature [containing] present in test data point [True]
21 Text feature [two] present in test data point [True]
23 Text feature [region] present in test data point [True]
24 Text feature [either] present in test data point [True]
25 Text feature [binding] present in test data point [True]
26 Text feature [function] present in test data point [True]
27 Text feature [results] present in test data point [True]
28 Text feature [expression] present in test data point [True]
29 Text feature [therefore] present in test data point [True]
30 Text feature [reduced] present in test data point [True]
31 Text feature [also] present in test data point [True]
33 Text feature [using] present in test data point [True]
34 Text feature [role] present in test data point [True]
35 Text feature [loss] present in test data point [True]
36 Text feature [shown] present in test data point [True]
37 Text feature [control] present in test data point [True]
38 Text feature [specific] present in test data point [True]
39 Text feature [table] present in test data point [True]
40 Text feature [affect] present in test data point [True]
41 Text feature [human] present in test data point [True]
42 Text feature [effect] present in test data point [True]
43 Text feature [gene] present in test data point [True]
44 Text feature [however] present in test data point [True]
45 Text feature [indicate] present in test data point [True]
46 Text feature [possible] present in test data point [True]
47 Text feature [following] present in test data point [True]
49 Text feature [four] present in test data point [True]
50 Text feature [determined] present in test data point [True]
```

```

51 Text feature [three] present in test data point [True]
53 Text feature [similar] present in test data point [True]
54 Text feature [important] present in test data point [True]
55 Text feature [critical] present in test data point [True]
56 Text feature [respectively] present in test data point [True]
60 Text feature [amino] present in test data point [True]
61 Text feature [well] present in test data point [True]
62 Text feature [previous] present in test data point [True]
65 Text feature [form] present in test data point [True]
66 Text feature [analysis] present in test data point [True]
67 Text feature [indicating] present in test data point [True]
68 Text feature [complex] present in test data point [True]
69 Text feature [indicated] present in test data point [True]
70 Text feature [proteins] present in test data point [True]
71 Text feature [previously] present in test data point [True]
73 Text feature [addition] present in test data point [True]
75 Text feature [corresponding] present in test data point [True]
77 Text feature [even] present in test data point [True]
78 Text feature [compared] present in test data point [True]
81 Text feature [result] present in test data point [True]
82 Text feature [least] present in test data point [True]
83 Text feature [25] present in test data point [True]
84 Text feature [used] present in test data point [True]
85 Text feature [including] present in test data point [True]
87 Text feature [acids] present in test data point [True]
88 Text feature [significant] present in test data point [True]
91 Text feature [may] present in test data point [True]
92 Text feature [within] present in test data point [True]
93 Text feature [essential] present in test data point [True]
94 Text feature [observed] present in test data point [True]
95 Text feature [furthermore] present in test data point [True]
97 Text feature [terminal] present in test data point [True]
Out of the top 100 features 65 are present in query point

```

0.4.3 K Nearest Neighbour Classification

```

In [69]: alpha = [5, 11, 15, 21, 31, 41, 51, 99]
         cv_log_error_array = []
         for i in alpha:
             print("for alpha =", i)
             clf = KNeighborsClassifier(n_neighbors=i)
             clf.fit(train_vec, train_y)
             sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
             sig_clf.fit(train_vec, train_y)
             sig_clf_probs = sig_clf.predict_proba(cv_vec)
             cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=
             # to avoid rounding error while multiplying probabilities we use log-probability e

```

```

        print("Log Loss :",log_loss(cv_y, sig_clf_probs))


fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],str(txt)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
clf.fit(train_vec, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_vec, train_y)

predict_y = sig_clf.predict_proba(train_vec)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log_
predict_y = sig_clf.predict_proba(test_vec)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss
predict_y = sig_clf.predict_proba(cv_vec)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",log_

for alpha = 5
Log Loss : 1.3265398215899902
for alpha = 11
Log Loss : 1.4042284081495355
for alpha = 15
Log Loss : 1.4409212647173233
for alpha = 21
Log Loss : 1.467403544557008
for alpha = 31
Log Loss : 1.5138769837737722
for alpha = 41
Log Loss : 1.5456574333802615
for alpha = 51
Log Loss : 1.5343982678526018
for alpha = 99
Log Loss : 1.4213936169591672

```



For values of best alpha = 5 The train log loss is: 1.020680894820197
 For values of best alpha = 5 The cross validation log loss is: 1.3290953937242058
 For values of best alpha = 5 The test log loss is: 1.3265398215899902


Testing the model with best hyper paramters

```
In [70]: clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
         predict_and_plot_confusion_matrix(train_vec, train_y, cv_vec, cv_y, clf)
```


Log loss : 1.3265398215899902

Number of mis-classified points : 0.4266917293233083


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



Feature Importance, Inorrectly classified point

```
In [77]: clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
         clf.fit(train_vec, train_y)
         sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
         sig_clf.fit(train_vec, train_y)

         test_point_index = 1

         predicted_cls = sig_clf.predict(test_vec[0].reshape(1,-1))
         print("Predicted Class :", predicted_cls[0])
         print("Actual Class :", test_y[test_point_index])
         neighbors = clf.kneighbors(test_vec[test_point_index].reshape(1, -1), alpha[best_alpha])
         print("The ", alpha[best_alpha], " nearest neighbours of the test points belongs to class")
         print("Frequency of nearest points :", Counter(train_y[neighbors[1][0]]))
```

Predicted Class : 4

Actual Class : 7

The 5 nearest neighbours of the test points belongs to classes [7 7 7 7 7]

Frequency of nearest points : Counter({7: 5})

Feature Importance, Correctly classified point

```
In [74]: clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
         clf.fit(train_vec, train_y)
         sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
         sig_clf.fit(train_vec, train_y)

         test_point_index = 1
```



```

predicted_cls = sig_clf.predict(test_vec[test_point_index].reshape(1,-1))
print("Predicted Class :", predicted_cls[0])
print("Actual Class :", test_y[test_point_index])
neighbors = clf.kneighbors(test_vec[test_point_index].reshape(1, -1), alpha[best_alpha])
print("the k value for knn is",alpha[best_alpha],"and the nearest neighbours of the test points belongs to classes [7 7]")
print("Fequency of nearest points :",Counter(train_y[neighbors[1][0]]))

```

Predicted Class : 7

Actual Class : 7

the k value for knn is 5 and the nearest neighbours of the test points belongs to classes [7 7]

Fequency of nearest points : Counter({7: 5})

0.4.4 Logistic Regression with class balancing

```

In [78]: alpha = [10 ** x for x in range(-6, 3)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2', loss='log', random_state=0)
    clf.fit(train_vec, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_vec, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_vec)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    # to avoid rounding error while multiplying probabilities we use log-probability
    print("Log Loss :",log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],str(txt)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=0)
clf.fit(train_vec, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_vec, train_y)

predict_y = sig_clf.predict_proba(train_vec)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log_loss(train_y, predict_y))
predict_y = sig_clf.predict_proba(cv_vec)

```

```


print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_loss)
predict_y = sig_clf.predict_proba(test_vec)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss)

```

```

for alpha = 1e-06
Log Loss : 1.3762267471928669
for alpha = 1e-05
Log Loss : 1.3858761287675025
for alpha = 0.0001
Log Loss : 1.3251788401733469
for alpha = 0.001
Log Loss : 1.0974391839339572
for alpha = 0.01
Log Loss : 1.098482224226122
for alpha = 0.1
Log Loss : 1.206796920687988
for alpha = 1
Log Loss : 1.4435460203183215
for alpha = 10
Log Loss : 1.486851536803223
for alpha = 100
Log Loss : 1.4914874092809136

```



```

For values of best alpha = 0.001 The train log loss is: 0.5611355366800372
For values of best alpha = 0.001 The cross validation log loss is: 1.0974391839339572

```

For values of best alpha = 0.001 The test log loss is: 1.0649284801253855

Testing the model with best hyper paramters

```
In [80]: ### clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2')
         predict_and_plot_confusion_matrix(train_vec, train_y, cv_vec, cv_y, clf)
```


Log loss : 1.0974391839339572

Number of mis-classified points : 0.34962406015037595

----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



Feature Importance, Correctly classified point

```
In [81]: # from tabulate import tabulate
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', l1_ratio=0.1)
clf.fit(train_vec, train_y)
test_point_index = 1
no_feature = 500
predicted_cls = sig_clf.predict(test_vec[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_vec[test_point_index]), 4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_[predicted_cls-1][:, :no_feature])
print("-"*50)
get_impfeature_names_tf(indices[0], test_df['TEXT'].iloc[test_point_index], test_df['G'])
```

Predicted Class : 7

Predicted Class Probabilities: [[0.0241 0.0467 0.0147 0.0246 0.0221 0.0141 0.8443 0.0048 0.0041]]

Actual Class : 7

```
-----
36 Text feature [constitutive] present in test data point [True]
65 Text feature [constitutively] present in test data point [True]
114 Text feature [activated] present in test data point [True]
158 Text feature [phosphotyrosine] present in test data point [True]
235 Text feature [activation] present in test data point [True]
371 Text feature [technology] present in test data point [True]
455 Text feature [interleukin] present in test data point [True]
488 Text feature [kinase] present in test data point [True]
Out of the top 500 features 8 are present in query point
```

Feature Importance, Incorrectly classified point

```
In [82]: test_point_index = 100
no_feature = 500
predicted_cls = sig_clf.predict(test_vec[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_vec[test_
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:, :no_feature]
print("-"*50)
get_impfeature_names_tf(indices[0], test_df['TEXT'].iloc[test_point_index], test_df['G
```

```
Predicted Class : 2
Predicted Class Probabilities: [[0.1653 0.2326 0.0324 0.1208 0.1394 0.0842 0.2039 0.0101 0.0111]
Actual Class : 6
```

```

116 Text feature [subgroup] present in test data point [True]
261 Text feature [median] present in test data point [True]
404 Text feature [formalin] present in test data point [True]
412 Text feature [therapy] present in test data point [True]
494 Text feature [ffpe] present in test data point [True]
Out of the top 500 features 5 are present in query point

```

0.4.5 Logistic Regression without class balancing

```
In [83]: alpha = [10 ** x for x in range(-6, 1)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_vec, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_vec, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_vec)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()
```


```

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_vec, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_vec, train_y)

predict_y = sig_clf.predict_proba(train_vec)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(train_vec, train_y, predict_y))
predict_y = sig_clf.predict_proba(cv_vec)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_loss(cv_vec, cv_y, predict_y))
predict_y = sig_clf.predict_proba(test_vec)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(test_vec, test_y, predict_y))

for alpha = 1e-06
Log Loss : 1.3587304074753803
for alpha = 1e-05
Log Loss : 1.3624834639880334
for alpha = 0.0001
Log Loss : 1.3225394674388362
for alpha = 0.001
Log Loss : 1.1135641054292025
for alpha = 0.01
Log Loss : 1.1294188383338022
for alpha = 0.1
Log Loss : 1.1662444045999807
for alpha = 1
Log Loss : 1.438973960607487

```



For values of best alpha = 0.001 The train log loss is: 0.5565790044926407
 For values of best alpha = 0.001 The cross validation log loss is: 1.1135641054292025
 For values of best alpha = 0.001 The test log loss is: 1.0729974789587016

Testing the model with best hyper paramters

```
In [87]: clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
         predict_and_plot_confusion_matrix(train_vec, train_y, cv_vec, cv_y, clf)
```


Log loss : 1.1135641054292025

Number of mis-classified points : 0.34398496240601506


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



Feature Importance, Correctly classified point

```
In [85]: clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
         clf.fit(train_vec, train_y)
         test_point_index = 1
         no_feature = 500
         predicted_cls = sig_clf.predict(test_vec[test_point_index])
         print("Predicted Class :", predicted_cls[0])
         print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_vec[test_point_index]), 4))
         print("Actual Class :", test_y[test_point_index])
         indices = np.argsort(-clf.coef_[predicted_cls-1][:, :no_feature])
         print("-"*50)
         get_impfeature_names_tf(indices[0], test_df['TEXT'].iloc[test_point_index], test_df['G'])
```

Predicted Class : 7

Predicted Class Probabilities: [[0.0292 0.0523 0.01 0.0309 0.0189 0.0136 0.8383 0.0046 0.0023]]

Actual Class : 7

```
-----
70 Text feature [constitutively] present in test data point [True]
101 Text feature [constitutive] present in test data point [True]
130 Text feature [activated] present in test data point [True]
200 Text feature [activation] present in test data point [True]
278 Text feature [phosphotyrosine] present in test data point [True]
299 Text feature [inhibited] present in test data point [True]
320 Text feature [technology] present in test data point [True]
376 Text feature [tyrosyl] present in test data point [True]
393 Text feature [kinase] present in test data point [True]
395 Text feature [phosphorylation] present in test data point [True]
413 Text feature [independence] present in test data point [True]
414 Text feature [ligand] present in test data point [True]
421 Text feature [py] present in test data point [True]
```

```

426 Text feature [interleukin] present in test data point [True]
439 Text feature [oncogene] present in test data point [True]
482 Text feature [expressing] present in test data point [True]
Out of the top 500 features 16 are present in query point

```

Feature Importance, Incorrectly classified point

```
In [86]: test_point_index = 100
no_feature = 500
predicted_cls = sig_clf.predict(test_vec[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_vec[test_
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:, :no_feature]
print("-"*50)
get_impfeature_names_tf(indices[0], test_df['TEXT'].iloc[test_point_index], test_df['G
```

```
Predicted Class : 2
Predicted Class Probabilities: [[0.1703 0.2399 0.0204 0.1178 0.1314 0.0821 0.223 0.0096 0.0055]]
Actual Class : 6
```

```

106 Text feature [subgroup] present in test data point [True]
268 Text feature [median] present in test data point [True]
326 Text feature [ffpe] present in test data point [True]
388 Text feature [therapy] present in test data point [True]
391 Text feature [formalin] present in test data point [True]
Out of the top 500 features 5 are present in query point

```

0.4.6 Random Forest Classifier

```
In [88]: alpha = [100,200,500,1000,2000]
max_depth = [5, 10]
cv_log_error_array = []
for i in alpha:
    for j in max_depth:
        print("for n_estimators =", i,"and max depth = ", j)
        clf = RandomForestClassifier(n_estimators=i, criterion='gini', max_depth=j, ra
        clf.fit(train_vec, train_y)
        sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
        sig_clf.fit(train_vec, train_y)
        sig_clf_probs = sig_clf.predict_proba(cv_vec)
        cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, e
        print("Log Loss :",log_loss(cv_y, sig_clf_probs))

'''fig, ax = plt.subplots()
features = np.dot(np.array(alpha)[: ,None],np.array(max_depth)[None]).ravel()
```

```

ax.plot(features, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[int(i/2)],max_depth[int(i%2)],str(txt)), (features[i],cv_log_e
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()
'''

best_alpha = np.argmin(cv_log_error_array)
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)], criterion='gini',
clf.fit(train_vec, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_vec, train_y)

predict_y = sig_clf.predict_proba(train_vec)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The train log loss")
predict_y = sig_clf.predict_proba(cv_vec)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The cross validation log loss")
predict_y = sig_clf.predict_proba(test_vec)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The test log loss")

for n_estimators = 100 and max depth = 5
Log Loss : 1.2417309452150282
for n_estimators = 100 and max depth = 10
Log Loss : 1.1912829351753556
for n_estimators = 200 and max depth = 5
Log Loss : 1.228500684679421
for n_estimators = 200 and max depth = 10
Log Loss : 1.1812957438294787
for n_estimators = 500 and max depth = 5
Log Loss : 1.2236113739683498
for n_estimators = 500 and max depth = 10
Log Loss : 1.1763219005393661
for n_estimators = 1000 and max depth = 5
Log Loss : 1.2277077844020792
for n_estimators = 1000 and max depth = 10
Log Loss : 1.1740521033604074
for n_estimators = 2000 and max depth = 5
Log Loss : 1.2279677047180988
for n_estimators = 2000 and max depth = 10
Log Loss : 1.1721777847644939
For values of best estimator = 2000 The train log loss is: 0.6786572369508065
For values of best estimator = 2000 The cross validation log loss is: 1.1721777847644936
For values of best estimator = 2000 The test log loss is: 1.1609443840222373

```


Testing the model with best hyper paramters

```
In [89]: clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)], criterion='gini',  
    predict_and_plot_confusion_matrix(train_vec, train_y,cv_vec,cv_y, clf)
```

Log loss : 1.1721777847644936

Number of mis-classified points : 0.37781954887218044


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



Feature Importance, Correctly classified point

```
In [90]: # test_point_index = 10
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)], criterion='gini',
clf.fit(train_vec, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_vec, train_y)

test_point_index = 1
no_feature = 100
predicted_cls = sig_clf.predict(test_vec[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_vec[test_
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.feature_importances_)
print("-"*50)
get_impfeature_names_tf(indices[:no_feature], test_df['TEXT'].iloc[test_point_index],
```

Predicted Class : 7

Predicted Class Probabilities: [[0.0294 0.1412 0.0157 0.0274 0.0374 0.0274 0.7136 0.0045 0.003

Actual Class : 7

```
-----
0 Text feature [activating] present in test data point [True]
1 Text feature [kinase] present in test data point [True]
2 Text feature [activation] present in test data point [True]
3 Text feature [tyrosine] present in test data point [True]
4 Text feature [phosphorylation] present in test data point [True]
5 Text feature [treatment] present in test data point [True]
6 Text feature [inhibitor] present in test data point [True]
7 Text feature [activated] present in test data point [True]
8 Text feature [oncogenic] present in test data point [True]
```

9 Text feature [inhibitors] present in test data point [True]
 10 Text feature [function] present in test data point [True]
 11 Text feature [therapy] present in test data point [True]
 12 Text feature [constitutive] present in test data point [True]
 13 Text feature [growth] present in test data point [True]
 15 Text feature [signaling] present in test data point [True]
 16 Text feature [drug] present in test data point [True]
 18 Text feature [constitutively] present in test data point [True]
 22 Text feature [trials] present in test data point [True]
 23 Text feature [resistance] present in test data point [True]
 24 Text feature [inhibition] present in test data point [True]
 25 Text feature [receptor] present in test data point [True]
 28 Text feature [treated] present in test data point [True]
 31 Text feature [months] present in test data point [True]
 32 Text feature [clinical] present in test data point [True]
 33 Text feature [protein] present in test data point [True]
 37 Text feature [kinases] present in test data point [True]
 40 Text feature [dose] present in test data point [True]
 41 Text feature [proliferation] present in test data point [True]
 45 Text feature [ic50] present in test data point [True]
 47 Text feature [amplification] present in test data point [True]
 48 Text feature [patients] present in test data point [True]
 54 Text feature [expressing] present in test data point [True]
 60 Text feature [phospho] present in test data point [True]
 61 Text feature [advanced] present in test data point [True]
 62 Text feature [potential] present in test data point [True]
 63 Text feature [oncogene] present in test data point [True]
 65 Text feature [cells] present in test data point [True]
 66 Text feature [resistant] present in test data point [True]
 68 Text feature [autophosphorylation] present in test data point [True]
 70 Text feature [nsccl] present in test data point [True]
 72 Text feature [independent] present in test data point [True]
 75 Text feature [ligand] present in test data point [True]
 77 Text feature [phosphorylated] present in test data point [True]
 81 Text feature [patient] present in test data point [True]
 82 Text feature [effective] present in test data point [True]
 84 Text feature [predicted] present in test data point [True]
 85 Text feature [sensitive] present in test data point [True]
 88 Text feature [inhibited] present in test data point [True]
 89 Text feature [active] present in test data point [True]
 94 Text feature [sensitivity] present in test data point [True]
 95 Text feature [response] present in test data point [True]
 96 Text feature [cell] present in test data point [True]
 97 Text feature [survival] present in test data point [True]
 99 Text feature [tki] present in test data point [True]
 Out of the top 100 features 54 are present in query point

Feature Importance, Incorrectly classified point

```
In [91]: test_point_index = 100
no_feature = 100
predicted_cls = sig_clf.predict(test_vec[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_vec[test_point_index])[0], 4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.feature_importances_)
print("-"*50)
get_impfeature_names_tf(indices[:no_feature], test_df['TEXT'].iloc[test_point_index], test_y[test_point_index])
```

Predicted Class : 2

Predicted Class Probabilities: `[[0.2245 0.2781 0.0249 0.1615 0.0702 0.0635 0.1603 0.0086 0.0086]]`

Actual Class : 6

```

1 Text feature [kinase] present in test data point [True]
3 Text feature [tyrosine] present in test data point [True]
5 Text feature [treatment] present in test data point [True]
6 Text feature [inhibitor] present in test data point [True]
9 Text feature [inhibitors] present in test data point [True]
10 Text feature [function] present in test data point [True]
11 Text feature [therapy] present in test data point [True]
13 Text feature [growth] present in test data point [True]
14 Text feature [suppressor] present in test data point [True]
25 Text feature [receptor] present in test data point [True]
26 Text feature [therapeutic] present in test data point [True]
32 Text feature [clinical] present in test data point [True]
33 Text feature [protein] present in test data point [True]
47 Text feature [amplification] present in test data point [True]
48 Text feature [patients] present in test data point [True]
49 Text feature [variants] present in test data point [True]
52 Text feature [functional] present in test data point [True]
54 Text feature [expressing] present in test data point [True]
56 Text feature [factor] present in test data point [True]
57 Text feature [harboring] present in test data point [True]
59 Text feature [stability] present in test data point [True]
62 Text feature [potential] present in test data point [True]
63 Text feature [oncogene] present in test data point [True]
65 Text feature [cells] present in test data point [True]
69 Text feature [brca1] present in test data point [True]
71 Text feature [functions] present in test data point [True]
72 Text feature [independent] present in test data point [True]
80 Text feature [variant] present in test data point [True]
81 Text feature [patient] present in test data point [True]
85 Text feature [sensitive] present in test data point [True]
90 Text feature [repair] present in test data point [True]
91 Text feature [pathway] present in test data point [True]
94 Text feature [sensitivity] present in test data point [True]

```

```

95 Text feature [response] present in test data point [True]
96 Text feature [cell] present in test data point [True]
97 Text feature [survival] present in test data point [True]
Out of the top 100 features 36 are present in query point

```

0.4.7 Stack the models

```

In [92]: clf1 = SGDClassifier(alpha=0.001, penalty='l2', loss='log', class_weight='balanced', random_state=42)
        clf1.fit(train_vec, train_y)
        sig_clf1 = CalibratedClassifierCV(clf1, method="sigmoid")

        clf2 = SGDClassifier(alpha=1, penalty='l2', loss='hinge', class_weight='balanced', random_state=42)
        clf2.fit(train_vec, train_y)
        sig_clf2 = CalibratedClassifierCV(clf2, method="sigmoid")

        clf3 = MultinomialNB(alpha=0.001)
        clf3.fit(train_vec, train_y)
        sig_clf3 = CalibratedClassifierCV(clf3, method="sigmoid")

        sig_clf1.fit(train_vec, train_y)
        print("Logistic Regression : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf1.predict_proba(cv_vec))))
        sig_clf2.fit(train_vec, train_y)
        print("Support vector machines : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf2.predict_proba(cv_vec))))
        sig_clf3.fit(train_vec, train_y)
        print("Naive Bayes : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf3.predict_proba(cv_vec))))
        print("-"*50)
        alpha = [0.0001,0.001,0.01,0.1,1,10]
        best_alpha = 999
        for i in alpha:
            lr = LogisticRegression(C=i)
            sclf = StackingClassifier(classifiers=[sig_clf1, sig_clf2, sig_clf3], meta_classifier=lr)
            sclf.fit(train_vec, train_y)
            print("Stacking Classifier : for the value of alpha: %f Log Loss: %0.3f" % (i, log_loss(cv_y, sclf.predict_proba(cv_vec))))
            log_error = log_loss(cv_y, sclf.predict_proba(cv_vec))
            if best_alpha > log_error:
                best_alpha = log_error

Logistic Regression : Log Loss: 1.11
Support vector machines : Log Loss: 1.49
Naive Bayes : Log Loss: 1.27
-----
Stacking Classifier : for the value of alpha: 0.000100 Log Loss: 2.178
Stacking Classifier : for the value of alpha: 0.001000 Log Loss: 2.038
Stacking Classifier : for the value of alpha: 0.010000 Log Loss: 1.516
Stacking Classifier : for the value of alpha: 0.100000 Log Loss: 1.130
Stacking Classifier : for the value of alpha: 1.000000 Log Loss: 1.218

```


Stacking Classifier : for the value of alpha: 10.000000 Log Loss: 1.464

Testing the model with best hyper paramters

```
In [93]: lr = LogisticRegression(C=0.1)
         sclf = StackingClassifier(classifiers=[sig_clf1, sig_clf2, sig_clf3], meta_classifier=lr)
         sclf.fit(train_vec, train_y)

         log_error = log_loss(train_y, sclf.predict_proba(train_vec))
         print("Log loss (train) on the stacking classifier :",log_error)

         log_error = log_loss(cv_y, sclf.predict_proba(cv_vec))
         print("Log loss (CV) on the stacking classifier :",log_error)

         log_error = log_loss(test_y, sclf.predict_proba(test_vec))
         print("Log loss (test) on the stacking classifier :",log_error)

         print("Number of missclassified point :", np.count_nonzero((sclf.predict(test_vec)- test_y)))
         plot_confusion_matrix(test_y=test_y, predict_y=sclf.predict(test_vec))
```


Log loss (train) on the stacking classifier : 0.6328410190333948

Log loss (CV) on the stacking classifier : 1.1296450689631101

Log loss (test) on the stacking classifier : 1.1091185225239562

Number of missclassified point : 0.3548872180451128


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



0.4.8 Maximum Voting classifier

```
In [94]: #Refer:http://scikit-learn.org/stable/modules/generated/sklearn.ensemble.VotingClassifier
from sklearn.ensemble import VotingClassifier
vclf = VotingClassifier(estimators=[('lr', sig_clf1), ('svc', sig_clf2), ('rf', sig_clf3)])
vclf.fit(train_vec, train_y)
print("Log loss (train) on the VotingClassifier :", log_loss(train_y, vclf.predict_proba(train_vec)))
print("Log loss (CV) on the VotingClassifier :", log_loss(cv_y, vclf.predict_proba(cv_vec)))
print("Log loss (test) on the VotingClassifier :", log_loss(test_y, vclf.predict_proba(test_vec)))
print("Number of missclassified point :", np.count_nonzero(vclf.predict(test_vec) - test_y))
plot_confusion_matrix(test_y=test_y, predict_y=vclf.predict(test_vec))
```


Log loss (train) on the VotingClassifier : 0.8565659723221687

Log loss (CV) on the VotingClassifier : 1.1576533848750443


Log loss (test) on the VotingClassifier : 1.1643755127972553

Number of missclassified point : 0.35789473684210527


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



0.5 Instead of using all the words in the dataset, use only the top 1000 words based of tf-idf values

```
In [15]: text_vectorizer.idf_.shape
```

```
Out[15]: (126099,)
```

```
In [16]: train_text_feature.shape
```

```
Out[16]: (2124, 126099)
```

```
In [17]: df=pd.DataFrame(text_vectorizer.get_feature_names())
```

```
In [18]: mean_tf=np.mean(train_text_feature,axis=0)
mean_tf
```

```
Out[18]: matrix([[0.00410664, 0.00943086, 0.00065508, ..., 0.00047081, 0.00047081,
0.00047081]])
```

```
In [19]: x=np.array(mean_tf)[0].tolist()
```

```
In [20]: important_feat=[]
important_feat=np.argsort((x))[::-1]
important_feat=important_feat[:1000]
important_feat.shape
```

```
Out[20]: (1000,)
```

```
In [21]: np.argmax(mean_tf)
```

```
Out[21]: 19795
```

```
In [22]: np.max(mean_tf)
```

```
Out [22]: 0.019041245324341777
```

```
In [23]: imp_feat=[]
        for index in important_feat:
            imp_feat.append(df.iloc[index])
```

```
In [24]: x=np.array(imp_feat)
        feature=[]
        for i in x:
            for j in i:
                feature.append(j)
```

```
In [25]: text_vectorizer=TfidfVectorizer()
        text_vectorizer.fit(feature)
        train=text_vectorizer.transform(train_df['TEXT'])
        test=text_vectorizer.transform(test_df['TEXT'])
        cv=text_vectorizer.transform(cv_df['TEXT'])
```

```
In [26]: train_vec = hstack((train_gene_feature_onehotCoding,train_variation_feature_onehotCoding))
        test_vec = hstack((test_gene_feature_onehotCoding,test_variation_feature_onehotCoding))
        cv_vec = hstack((cv_gene_feature_onehotCoding,cv_variation_feature_onehotCoding,cv))
```

0.5.1 Logistic Regression with class Balance

```
In [132]: alpha = [10 ** x for x in range(-6, 3)]
        cv_log_error_array = []
        for i in alpha:
            print("for alpha =", i)
            clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2', loss='log', )
            clf.fit(train_vec, train_y)
            sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
            sig_clf.fit(train_vec, train_y)
            sig_clf_probs = sig_clf.predict_proba(cv_vec)
            cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
            # to avoid rounding error while multiplying probabilities we use log-probability
            print("Log Loss :",log_loss(cv_y, sig_clf_probs))

        fig, ax = plt.subplots()
        ax.plot(alpha, cv_log_error_array,c='g')
        for i, txt in enumerate(np.round(cv_log_error_array,3)):
            ax.annotate((alpha[i],str(txt)), (alpha[i],cv_log_error_array[i]))
        plt.grid()
        plt.title("Cross Validation Error for each alpha")
        plt.xlabel("Alpha i's")
        plt.ylabel("Error measure")
        plt.show()

        best_alpha = np.argmin(cv_log_error_array)
```


```

clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2',
clf.fit(train_vec, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_vec, train_y)

predict_y = sig_clf.predict_proba(train_vec)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log_
predict_y = sig_clf.predict_proba(cv_vec)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log lo
predict_y = sig_clf.predict_proba(test_vec)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",log_

for alpha = 1e-06
Log Loss : 1.1011976653190054
for alpha = 1e-05
Log Loss : 1.0525242861483708
for alpha = 0.0001
Log Loss : 0.9932250639803677
for alpha = 0.001
Log Loss : 1.0253601459328905
for alpha = 0.01
Log Loss : 1.1704809031082442
for alpha = 0.1
Log Loss : 1.4806280813212045
for alpha = 1
Log Loss : 1.6569474163813263
for alpha = 10
Log Loss : 1.6786255997965651
for alpha = 100
Log Loss : 1.680908657292085

```



For values of best alpha = 0.0001 The train log loss is: 0.4627039169136912
 For values of best alpha = 0.0001 The cross validation log loss is: 0.9932250639803677
 For values of best alpha = 0.0001 The test log loss is: 1.0185008365416044

Testing the model with best hyper paramters

```
In [133]: clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2',
predict_and_plot_confusion_matrix(train_vec, train_y, cv_vec, cv_y, clf)
```


Log loss : 0.9932250639803677

Number of mis-classified points : 0.32706766917293234


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



```
In [113]: text_vectorizer
```

```
Out[113]: (2124, 1000)
```

```
In [137]: test_vec=test_vec.tocsr()
```

```
In [27]: #Feature importance for tfidf with 1000 features
```

```
def get_impfeature_names_tf_imp(indices, text, gene, var, no_features):
    gene_count_vec = CountVectorizer()
    var_count_vec = CountVectorizer()
```

```
gene_vec = gene_count_vec.fit(train_df['Gene'])
var_vec = var_count_vec.fit(train_df['Variation'])
```

```
fea1_len = len(gene_vec.get_feature_names())
fea2_len = len(var_count_vec.get_feature_names())
```

```
word_present = 0
```

```
for i,v in enumerate(indices):
```

```
    if (v < fea1_len):
```

```
        word = gene_vec.get_feature_names()[v]
```

```
        yes_no = True if word == gene else False
```

```
        if yes_no:
```

```
            word_present += 1
```

```
            print(i, "Gene feature [{}] present in test data point [{}]" .format(w
```

```
    elif (v < fea1_len+fea2_len):
```

```
        word = var_vec.get_feature_names()[v-(fea1_len)]
```

```
        yes_no = True if word == var else False
```

```
        if yes_no:
```

```
            word_present += 1
```

```

        print(i, "variation feature [{}]" + "present in test data point [{}]" + "format(w
    else:
        word = text_vectorizer.get_feature_names()[v-(fea1_len+fea2_len)]
        yes_no = True if word in text.split() else False
        if yes_no:
            word_present += 1
            print(i, "Text feature [{}]" + "present in test data point [{}]" + "format(w

    print("Out of the top ", no_features, " features ", word_present, "are present in q

```

Feature Importance, Correctly classified point

```

In [139]: # from tabulate import tabulate
          clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2',
          clf.fit(train_vec, train_y)
          test_point_index = 1
          no_feature = 500
          predicted_cls = sig_clf.predict(test_vec[test_point_index])
          print("Predicted Class :", predicted_cls[0])
          print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_vec[test_
          print("Actual Class :", test_y[test_point_index])
          indices = np.argsort(-clf.coef_)[predicted_cls-1][:, :no_feature]
          print("-"*50)
          get_impfeature_names_tf_imp(indices[0], test_df['TEXT'].iloc[test_point_index], test_

```

Predicted Class : 7

Predicted Class Probabilities: [[0.0204 0.1397 0.0162 0.0211 0.0532 0.0178 0.7217 0.0065 0.003

Actual Class : 7

```

-----
0 Text feature [activation] present in test data point [True]
17 Text feature [mutant] present in test data point [True]
18 Text feature [oncogenic] present in test data point [True]
20 Text feature [high] present in test data point [True]
21 Text feature [cells] present in test data point [True]
22 Text feature [tyrosine] present in test data point [True]
26 Text feature [cancers] present in test data point [True]
28 Text feature [signaling] present in test data point [True]
58 Text feature [fig] present in test data point [True]
65 Text feature [12] present in test data point [True]
71 Text feature [increased] present in test data point [True]
78 Text feature [fold] present in test data point [True]
90 Text feature [proliferation] present in test data point [True]
93 Text feature [per] present in test data point [True]
98 Text feature [also] present in test data point [True]
113 Text feature [previously] present in test data point [True]
135 Text feature [activated] present in test data point [True]
149 Text feature [mutants] present in test data point [True]
158 Text feature [total] present in test data point [True]

```

```

163 Text feature [wt] present in test data point [True]
219 Text feature [inhibitor] present in test data point [True]
224 Text feature [genomic] present in test data point [True]
255 Text feature [expressing] present in test data point [True]
261 Text feature [common] present in test data point [True]
285 Text feature [positive] present in test data point [True]
288 Text feature [presence] present in test data point [True]
292 Text feature [obtained] present in test data point [True]
306 Text feature [reported] present in test data point [True]
317 Text feature [found] present in test data point [True]
325 Text feature [two] present in test data point [True]
345 Text feature [kinase] present in test data point [True]
378 Text feature [10] present in test data point [True]
388 Text feature [14] present in test data point [True]
390 Text feature [cdna] present in test data point [True]
394 Text feature [could] present in test data point [True]
397 Text feature [survival] present in test data point [True]
400 Text feature [tissue] present in test data point [True]
404 Text feature [inhibited] present in test data point [True]
406 Text feature [tumor] present in test data point [True]
429 Text feature [000] present in test data point [True]
434 Text feature [patients] present in test data point [True]
452 Text feature [constitutively] present in test data point [True]
465 Text feature [18] present in test data point [True]
483 Text feature [alterations] present in test data point [True]
491 Text feature [deletion] present in test data point [True]
496 Text feature [50] present in test data point [True]
498 Text feature [approximately] present in test data point [True]
Out of the top 500 features 47 are present in query point

```

Feature Importance, Incorrectly classified point

```

In [142]: test_point_index = 300
          no_feature = 500
          predicted_cls = sig_clf.predict(test_vec[test_point_index])
          print("Predicted Class :", predicted_cls[0])
          print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_vec[test_point_index]), 5))
          print("Actual Class :", test_y[test_point_index])
          indices = np.argsort(-clf.coef_)[predicted_cls-1][:, :no_feature]
          print("-"*50)
          get_impfeature_names_tf_imp(indices[0], test_df['TEXT'].iloc[test_point_index], test_y[test_point_index])

Predicted Class : 7
Predicted Class Probabilities: [[0.0321 0.1555 0.0147 0.0322 0.0776 0.0115 0.6683 0.005 0.003]
Actual Class : 2
-----
0 Text feature [activation] present in test data point [True]

```

12 Text feature [somatic] present in test data point [True]
13 Text feature [downstream] present in test data point [True]
17 Text feature [mutant] present in test data point [True]
20 Text feature [high] present in test data point [True]
21 Text feature [cells] present in test data point [True]
22 Text feature [tyrosine] present in test data point [True]
28 Text feature [signaling] present in test data point [True]
58 Text feature [fig] present in test data point [True]
65 Text feature [12] present in test data point [True]
71 Text feature [increased] present in test data point [True]
78 Text feature [fold] present in test data point [True]
90 Text feature [proliferation] present in test data point [True]
98 Text feature [also] present in test data point [True]
99 Text feature [pathways] present in test data point [True]
113 Text feature [previously] present in test data point [True]
121 Text feature [factor] present in test data point [True]
135 Text feature [activated] present in test data point [True]
149 Text feature [mutants] present in test data point [True]
158 Text feature [total] present in test data point [True]
163 Text feature [wt] present in test data point [True]
182 Text feature [levels] present in test data point [True]
189 Text feature [position] present in test data point [True]
224 Text feature [genomic] present in test data point [True]
255 Text feature [expressing] present in test data point [True]
261 Text feature [common] present in test data point [True]
287 Text feature [2a] present in test data point [True]
288 Text feature [presence] present in test data point [True]
292 Text feature [obtained] present in test data point [True]
295 Text feature [transcription] present in test data point [True]
306 Text feature [reported] present in test data point [True]
312 Text feature [tissues] present in test data point [True]
317 Text feature [found] present in test data point [True]
325 Text feature [two] present in test data point [True]
345 Text feature [kinase] present in test data point [True]
378 Text feature [10] present in test data point [True]
388 Text feature [14] present in test data point [True]
390 Text feature [cdna] present in test data point [True]
394 Text feature [could] present in test data point [True]
397 Text feature [survival] present in test data point [True]
400 Text feature [tissue] present in test data point [True]
404 Text feature [inhibited] present in test data point [True]
406 Text feature [tumor] present in test data point [True]
425 Text feature [coding] present in test data point [True]
434 Text feature [patients] present in test data point [True]
445 Text feature [factors] present in test data point [True]
448 Text feature [activate] present in test data point [True]
452 Text feature [constitutively] present in test data point [True]
463 Text feature [without] present in test data point [True]

```

465 Text feature [18] present in test data point [True]
470 Text feature [higher] present in test data point [True]
479 Text feature [product] present in test data point [True]
483 Text feature [alterations] present in test data point [True]
493 Text feature [transformation] present in test data point [True]
496 Text feature [50] present in test data point [True]
498 Text feature [approximately] present in test data point [True]
Out of the top 500 features 56 are present in query point

```

```
In [28]: test_vec=test_vec.tocsr()
```

1 Apply Logistic regression with CountVectorizer Features, including both unigrams and bigrams

```
In [143]: text_vectorizer = CountVectorizer(min_df=3,ngram_range=(1,2))
         train_text_feature_onehotCoding = text_vectorizer.fit_transform(train_df['TEXT'])
         # getting all the feature names (words)
         train_text_features= text_vectorizer.get_feature_names()
```

```
print("Total number of unique words in train data :", len(train_text_features))
```

Total number of unique words in train data : 796397

```
In [144]: # don't forget to normalize every feature
         train_text_feature_onehotCoding = normalize(train_text_feature_onehotCoding, axis=0)

         # we use the same vectorizer that was trained on train data
         test_text_feature_onehotCoding = text_vectorizer.transform(test_df['TEXT'])
         # don't forget to normalize every feature
         test_text_feature_onehotCoding = normalize(test_text_feature_onehotCoding, axis=0)
```

```
         # we use the same vectorizer that was trained on train data
         cv_text_feature_onehotCoding = text_vectorizer.transform(cv_df['TEXT'])
         # don't forget to normalize every feature
         cv_text_feature_onehotCoding = normalize(cv_text_feature_onehotCoding, axis=0)
```

```
In [145]: train_vec = hstack((train_gene_feature_onehotCoding,train_variation_feature_onehotCoding))
         test_vec = hstack((test_gene_feature_onehotCoding,test_variation_feature_onehotCoding))
         cv_vec = hstack((cv_gene_feature_onehotCoding,cv_variation_feature_onehotCoding,cv_text_feature_onehotCoding))
```

1.0.1 Logistic regression

```
In [12]: #feature importance for 2gram BoW
         def get_impfeature_names_bow_2gram(indices, text, gene, var, no_features):
             gene_count_vec = CountVectorizer()
```

```

var_count_vec = CountVectorizer()
text_count_vec = CountVectorizer(min_df=3,ngram_range=(1,2))

gene_vec = gene_count_vec.fit(train_df['Gene'])
var_vec = var_count_vec.fit(train_df['Variation'])
text_vec = text_count_vec.fit(train_df['TEXT'])

fea1_len = len(gene_vec.get_feature_names())
fea2_len = len(var_count_vec.get_feature_names())

word_present = 0
for i,v in enumerate(indices):
    if (v < fea1_len):
        word = gene_vec.get_feature_names()[v]
        yes_no = True if word == gene else False
        if yes_no:
            word_present += 1
            print(i, "Gene feature [{}] present in test data point [{}]" .format(w
    elif (v < fea1_len+fea2_len):
        word = var_vec.get_feature_names()[v-(fea1_len)]
        yes_no = True if word == var else False
        if yes_no:
            word_present += 1
            print(i, "variation feature [{}] present in test data point [{}]" .form
    else:
        word = text_vec.get_feature_names()[v-(fea1_len+fea2_len)]
        yes_no = True if word in text.split() else False
        if yes_no:
            word_present += 1
            print(i, "Text feature [{}] present in test data point [{}]" .format(w

print("Out of the top ",no_features," features ", word_present, "are present in q

```

```

In [146]: alpha = [10 ** x for x in range(-6, 3)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2', loss='log',
    clf.fit(train_vec, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_vec, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_vec)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps
    # to avoid rounding error while multiplying probabilities we use log-probability
    print("Log Loss :",log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array,c='g')

```

```


for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],str(txt)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', )
clf.fit(train_vec, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_vec, train_y)

predict_y = sig_clf.predict_proba(train_vec)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log_
predict_y = sig_clf.predict_proba(cv_vec)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:",log_
predict_y = sig_clf.predict_proba(test_vec)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",log_)

for alpha = 1e-06
Log Loss : 1.5452224097324696
for alpha = 1e-05
Log Loss : 1.559690091434177
for alpha = 0.0001
Log Loss : 1.5438538989527038
for alpha = 0.001
Log Loss : 1.4707385242692435
for alpha = 0.01
Log Loss : 1.1566704885794588
for alpha = 0.1
Log Loss : 1.1744047319616486
for alpha = 1
Log Loss : 1.2573623352622487
for alpha = 10
Log Loss : 1.3337836693304905
for alpha = 100
Log Loss : 1.3498408414671528

```



For values of best alpha = 0.01 The train log loss is: 0.8472646352615477
 For values of best alpha = 0.01 The cross validation log loss is: 1.1566704885794588
 For values of best alpha = 0.01 The test log loss is: 1.176842431113536


Testing the model with best hyper paramters

```
In [147]: clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2',
                             predict_and_plot_confusion_matrix(train_vec, train_y, cv_vec, cv_y, clf)
```


Log loss : 1.1566704885794588

Number of mis-classified points : 0.39097744360902253


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



```
In [227]: text_count_vec = CountVectorizer(min_df=3,ngram_range=(1,2))
          text_vec = text_count_vec.fit(train_df['TEXT'])
```

feature Importance, Correctly classified point

```
In [150]: test_vec=test_vec.tocsr()
```

```
In [156]: # from tabulate import tabulate
          clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2',
          clf.fit(train_vec,train_y)
          test_point_index = 1
          no_feature = 1000
          predicted_cls = sig_clf.predict(test_vec[test_point_index])
          print("Predicted Class :", predicted_cls[0])
          print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_vec[test_
          print("Actual Class :", test_y[test_point_index])
          indices = np.argsort(-clf.coef_)[predicted_cls-1][:,:no_feature]
          print("-"*50)
          get_impfeature_names_bow_2gram(indices[0], test_df['TEXT'].iloc[test_point_index],tes
```

Predicted Class : 7

Predicted Class Probabilities: [[0.0698 0.0799 0.015 0.0559 0.0225 0.0274 0.7194 0.0048 0.0055]

Actual Class : 7

882 Text feature [constitutively] present in test data point [True]
 Out of the top 1000 features 1 are present in query point

Feature Importance, Incorrectly classified point

```
In [157]: test_point_index = 100
          no_feature = 1000
```

```

predicted_cls = sig_clf.predict(test_vec[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_vec[test.
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,no_feature]
print("-"*50)
get_impfeature_names_bow_2gram(indices[0], test_df['TEXT'].iloc[test_point_index], te

```

Predicted Class : 7

Predicted Class Probabilities: [[0.1959 0.1817 0.0238 0.1397 0.0762 0.0701 0.2992 0.006 0.007

Actual Class : 6

Out of the top 1000 features 0 are present in query point

1.1 Feature Engineering

In [158]: train_df.head()

```

Out[158]:
      ID  Gene      Variation  Class  \
2426  2426  BRCA1          C64G      4
636   636  CDKN1A  Truncating Mutations      1
703   703  ERBB2          E812K      6
1455  1455  FGFR2          K659N      7
103   103  MSH6          R976H      1

      TEXT
2426  published analyses effects missense mutations ...
636   introduction loss control mammalian cell cycle...
703   purpose mutations associated resistance kinase...
1455  activating mutations tyrosine kinase domain re...
103   msh6 gene one mismatch repair genes involved h...

```

1.1.1 Text word length, character length and variation length as features

```

In [10]: train_df['word_len']=train_df['TEXT'].apply(lambda x : len(x.split()))
train_df['char_len']=train_df['TEXT'].apply(lambda x : len(x))
test_df['word_len']=test_df['TEXT'].apply(lambda x : len(x.split()))
test_df['char_len']=test_df['TEXT'].apply(lambda x : len(x))
cv_df['word_len']=cv_df['TEXT'].apply(lambda x : len(x.split()))
cv_df['char_len']=cv_df['TEXT'].apply(lambda x : len(x))
train_df.head()

```

```

Out[10]:
      ID  Gene      Variation  Class  \
2426  2426  BRCA1          C64G      4
636   636  CDKN1A  Truncating Mutations      1
703   703  ERBB2          E812K      6
1455  1455  FGFR2          K659N      7
103   103  MSH6          R976H      1

```

		TEXT	word_len	char_len
2426	published analyses effects missense mutations ...		17926	130265
636	introduction loss control mammalian cell cycle...		7234	53047
703	purpose mutations associated resistance kinase...		4961	37512
1455	activating mutations tyrosine kinase domain re...		8077	59908
103	msh6 gene one mismatch repair genes involved h...		13352	100388

```
In [14]: word_len_train=np.reshape(np.array(train_df['word_len']), (train_df['word_len'].shape[0],))
char_len_train=np.reshape(np.array(train_df['char_len']), (train_df['char_len'].shape[0],))
word_len_test=np.reshape(np.array(test_df['word_len']), (test_df['word_len'].shape[0],))
char_len_test=np.reshape(np.array(test_df['char_len']), (test_df['char_len'].shape[0],))
word_len_cv=np.reshape(np.array(cv_df['word_len']), (cv_df['word_len'].shape[0],))
char_len_cv=np.reshape(np.array(cv_df['char_len']), (cv_df['char_len'].shape[0],))
```

```
In [170]: train_df['Variation_len']=train_df['Variation'].apply(lambda x: len(x))
test_df['Variation_len']=test_df['Variation'].apply(lambda x: len(x))
cv_df['Variation_len']=cv_df['Variation'].apply(lambda x: len(x))
```

```
In [171]: train_df.head()
```

```
Out[171]:
```


	ID	Gene	Variation	Class	\
2426	2426	BRCA1	C64G	4	
636	636	CDKN1A	Truncating Mutations	1	
703	703	ERBB2	E812K	6	
1455	1455	FGFR2	K659N	7	
103	103	MSH6	R976H	1	

		TEXT	word_len	char_len	\
2426	published analyses effects missense mutations ...		17926	130265	
636	introduction loss control mammalian cell cycle...		7234	53047	
703	purpose mutations associated resistance kinase...		4961	37512	
1455	activating mutations tyrosine kinase domain re...		8077	59908	
103	msh6 gene one mismatch repair genes involved h...		13352	100388	

	Variation_len
2426	4
636	20
703	5
1455	5
103	5

```
In [244]: fig = plt.figure(figsize=(20,15))
sns.pairplot(train_df, hue='Class', vars=["Variation_len", 'word_len', 'char_len'])
plt.show()
```

<Figure size 1440x1080 with 0 Axes>



we can see that the features are not helping in classification and also adding these features did not improve the models.

1.1.2 Truncated svd to reduce the dimensions

```
In [207]: from sklearn.feature_extraction.text import TfidfVectorizer
text_vectorizer = TfidfVectorizer(max_features=10000)
train_text_feature = text_vectorizer.fit_transform(train_df['TEXT'])
# getting all the feature names (words)
train_text_features= text_vectorizer.get_feature_names()
test_text_feature = text_vectorizer.transform(test_df['TEXT'])
cv_text_feature = text_vectorizer.transform(cv_df['TEXT'])
```

```
In [208]: train_text_feature = normalize(train_text_feature, axis=0)
```

```
# don't forget to normalize every feature
test_text_feature = normalize(test_text_feature, axis=0)
```

```
cv_text_feature = normalize(cv_text_feature, axis=0)
```

```
In [211]: #stacking the features
```

```
train_vec = hstack((train_gene_feature_onehotCoding,train_variation_feature_onehotCoding))
test_vec = hstack((test_gene_feature_onehotCoding,test_variation_feature_onehotCoding))
cv_vec = hstack((cv_gene_feature_onehotCoding,cv_variation_feature_onehotCoding,cv_text_feature))
```

```
In [213]: #using tsvd to reduce the dimensions
```

```
from sklearn.decomposition import TruncatedSVD
```

```
from tqdm import tqdm
```

```
#applying truncated svd with n_components =1000
```

```
tsvd =TruncatedSVD(algorithm="randomized",n_components=2000, n_iter=7,\
                    random_state=42,tol=0.0)
```

```
X_reduced = tsvd.fit_transform(train_vec)
```

```
#plotting the percentage of the variance explained by the features
```

```
percentage_variance_explained=tsvd.explained_variance_/sum(tsvd.explained_variance_)
```

```
cum_variance_explained=np.cumsum(percentage_variance_explained)
```


```
plt.plot(cum_variance_explained,linewidth='2')
```

```
plt.xlabel("number of features")
```

```
plt.ylabel("variance")
```

```
plt.title("Percentage of variance explained")
```

```
plt.show()
```



```
In [220]: tsvd = TruncatedSVD(algorithm="randomized", n_components=1500, n_iter=7,\
                             random_state=42, tol=0.0)
tra = tsvd.fit_transform(train_vec)
tes = tsvd.transform(test_vec)
c = tsvd.transform(cv_vec)
```

1.2 Logistic Regression with Balancing

```
In [221]: alpha = [10 ** x for x in range(-6, 3)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2', loss='log',
                        random_state=42)
    clf.fit(tra, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(tra, train_y)
    sig_clf_probs = sig_clf.predict_proba(c)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    # to avoid rounding error while multiplying probabilities we use log-probability
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2',
                    random_state=42)
clf.fit(tra, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(tra, train_y)


predict_y = sig_clf.predict_proba(tra)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(tra, predict_y, labels=clf.classes_))
predict_y = sig_clf.predict_proba(c)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_loss(c, predict_y, labels=clf.classes_))
predict_y = sig_clf.predict_proba(tes)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(tes, predict_y, labels=clf.classes_))

for alpha = 1e-06
Log Loss : 1.0570658874759329
for alpha = 1e-05
```

```

Log Loss : 1.0465970705661962
for alpha = 0.0001
Log Loss : 0.9864306000414919
for alpha = 0.001
Log Loss : 0.9266671112993817
for alpha = 0.01
Log Loss : 1.0057874102939246
for alpha = 0.1
Log Loss : 1.0717122327927189
for alpha = 1
Log Loss : 1.267646051437711
for alpha = 10
Log Loss : 1.6622461851007915
for alpha = 100
Log Loss : 1.7372144829297962

```



```

For values of best alpha = 0.001 The train log loss is: 0.5726296501912288
For values of best alpha = 0.001 The cross validation log loss is: 0.9266671112993817
For values of best alpha = 0.001 The test log loss is: 0.9706599524824131

```

```


In [222]: clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2',
predict_and_plot_confusion_matrix(tra, train_y, c, cv_y, clf)

```


Log loss : 0.9266671112993817

Number of mis-classified points : 0.31954887218045114

----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



1.2.1 Logistic Regression without Class Balance

```
In [239]: alpha = [10 ** x for x in range(-6, 3)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(tra, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(tra, train_y)
    sig_clf_probs = sig_clf.predict_proba(c)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=
    # to avoid rounding error while multiplying probabilities we use log-probability
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=
clf.fit(tra, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(tra, train_y)
```

```


predict_y = sig_clf.predict_proba(tra)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log_
predict_y = sig_clf.predict_proba(c)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log lo
predict_y = sig_clf.predict_proba(tes)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",log_

```

```

for alpha = 1e-06
Log Loss : 1.0376463311304114
for alpha = 1e-05
Log Loss : 1.0288008319989888
for alpha = 0.0001
Log Loss : 0.9930706350433093
for alpha = 0.001
Log Loss : 0.9241393237294904
for alpha = 0.01
Log Loss : 0.9950353774753347
for alpha = 0.1
Log Loss : 1.0832774813439614
for alpha = 1
Log Loss : 1.241828571935979
for alpha = 10
Log Loss : 1.6463772476034435
for alpha = 100
Log Loss : 1.7628574422286374

```




For values of best alpha = 0.001 The train log loss is: 0.5605561048509772
 For values of best alpha = 0.001 The cross validation log loss is: 0.9241393237294904
 For values of best alpha = 0.001 The test log loss is: 0.9642156482902127

```
In [240]: clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=
          predict_and_plot_confusion_matrix(tra, train_y, c, cv_y, clf)
```


Log loss : 0.9241393237294904

Number of mis-classified points : 0.3233082706766917


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



1.3 Conclusion

Objective: Classify the given genetic variations/mutations based on evidence from text-based clinical literature.

1. We are given a dataset with 3321 rows and 5 columns: ID , Gene, Variations, Text, Class of which Class is a multiclass column containing the 9 classes.
2. We preprocess the data and split into train, test and cv set.
3. We perform EDA on the dataset to know about the distribution of the data and get more insights about the data and featurize the data using responsecoding and onehotencoding method.
4. We then check if distribution is stable across train test and cv set and then check how important gene and variation feature is in classifying the model.
5. We then apply various model like logistic Regression, Naive Bayes, Random Forest, Voting classifier etc and compare them.
6. We use multiclass log loss as error metric, confusion matrix as performance metric and also check the number of misclassified points.
7. Since interpretability is important we used only models which are interpretable and got the feature importance.
8. We then try to reduce the log loss using other techniques like tfidf vectorization, BoW Vectorization with 2 gram and tsvd.

```
In [241]: from prettytable import PrettyTable
```

```
x=PrettyTable()
```

```

x.field_names=['Algorithm','Vectorizer','penalty','alpha','CV Log loss','Test Log Loss']
x.add_row(["Naive Bayes","TFIDF","-",1,1.188, 1.176,39.47])
x.add_row(["KNN","TFIDF","-",5,1.329 ,1.326, 42.66 ])
x.add_row(["Logistic Regression (Balanced) ","TFIDF",'12',0.001,1.097, 1.064, 34.96 ])
x.add_row(["Logistic Regression","TFIDF",'12',0.001,1.113, 1.072 ,34.39 ])
x.add_row(["Random Forest","TFIDF","-",2000,1.172, 1.16, 37.78 ])
x.add_row(["Stacking","TFIDF","-","-",1.129 ,1.109, 35.48 ])
x.add_row(["Logistic Regression (Balanced)","TFIDF-1000",'12',0.0001,0.9932, 1.018 ,34.39 ])
x.add_row(["Logistic Regression (Balanced)","BoW 2 gram",'12',0.01,1.156 ,1.176, 39.47 ])
x.add_row(["Logistic Regression (Balanced)","TFIDF-TSVD",'12',0.001,0.9266, 0.9706, 32.33 ])
x.add_row(["Logistic Regression","TFIDF-TSVD",'12',0.001,0.924, 0.964, 32.33])

```

```
print(x)
```

Algorithm	Vectorizer	penalty	alpha	CV Log loss	Test Log Loss
Naive Bayes	TFIDF	-	1	1.188	1.176
KNN	TFIDF	-	5	1.329	1.326
Logistic Regression (Balanced)	TFIDF	12	0.001	1.097	1.064
Logistic Regression	TFIDF	12	0.001	1.113	1.072
Random Forest	TFIDF	-	2000	1.172	1.16
Stacking	TFIDF	-	-	1.129	1.109
Logistic Regression (Balanced)	TFIDF-1000	12	0.0001	0.9932	1.018
Logistic Regression (Balanced)	BoW 2 gram	12	0.01	1.156	1.176
Logistic Regression (Balanced)	TFIDF-TSVD	12	0.001	0.9266	0.9706
Logistic Regression	TFIDF-TSVD	12	0.001	0.924	0.964

Logistic Regression with TFIDf using tsvd reduced features has the log loss of 0.97 and lowest percent of misclassified points .

```
In [ ]:
```